

# hw02\_student\_solutions

September 20, 2019

## 1 Homework 2: Arrays and Tables

```
[ ]: # Don't change this cell; just run it.  
import numpy as np  
from datascience import *
```

**Recommended Reading:** \* [Data Types](#) \* [Sequences](#) \* [Tables](#)

Please complete this notebook by filling in the cells provided. Before you begin, execute the following cell to load the provided tests. Each time you start your server, you will need to execute this cell again to load the tests.

This assignment is due Thursday, September 12 at 11:59 P.M. You will receive an early submission bonus point if you turn in your final submission by Wednesday, September 11 at 11:59 P.M. Late work will not be accepted as per the [policies](#) of this course.

Start early so that you can come to office hours if you're stuck. Check the website for the office hours schedule.

Throughout this homework and all future ones, please be sure to not re-assign variables throughout the notebook! For example, if you use `max_temperature` in your answer to one question, do not reassign it later on.

**Important:** In this homework, the `ok` tests will tell you whether your answer is correct, except for Parts 4, 5 & 6. In future homework assignments, correctness tests will typically not be provided.

Before continuing the assignment, select "Save and Checkpoint" in the File menu and then execute the submit cell below. The result will contain a link that you can use to check that your assignment has been submitted successfully. If you submit more than once before the deadline, we will only grade your final submission. If you mistakenly submit the wrong one, you can head to [okpy.org](http://okpy.org) and flag the correct version. There will be another submit cell at the end of the assignment when you finish!

```
[ ]: # When you log-in please hit return (not shift + return) after typing in your  
      ↪ email  
_ = ok.submit()
```

## 1.1 1. Creating Arrays

**Question 1.** Make an array called `weird_numbers` containing the following numbers (in the given order):

1. -2
2. the sine of 1.2
3. 3
4. 5 to the power of the cosine of 1.2

*Hint:* `sin` and `cos` are functions in the `math` module.

BEGIN QUESTION

name: q1\_1

```
[2]: # Our solution involved one extra line of code before creating
      # weird_numbers.
      import math #SOLUTION
      weird_numbers = make_array(-2, math.sin(1.2), 3, 5**math.cos(1.2)) #SOLUTION
      weird_numbers
```

```
[2]: array([-2.          ,  0.93203909,  3.          ,  1.79174913])
```

```
[3]: # TEST
      np.allclose(weird_numbers, np.array([-2.,  0.93203909,  3.,  1.79174913]),
      ↪rtol=1e-03, atol=1e-03)
```

```
[3]: True
```

**Question 2.** Make an array called `book_title_words` containing the following three strings: "Eats", "Shoots", and "and Leaves".

BEGIN QUESTION

name: q1\_2

```
[4]: book_title_words = make_array("Eats", "Shoots", "and Leaves") #SOLUTION
      book_title_words
```

```
[4]: array(['Eats', 'Shoots', 'and Leaves'], dtype='<U10')
```

Strings have a method called `join`. `join` takes one argument, an array of strings. It returns a single string. Specifically, the value of `a_string.join(an_array)` is a single string that's the **concatenation** ("putting together") of all the strings in `an_array`, **except** `a_string` is inserted in between each string.

**Question 3.** Use the array `book_title_words` and the method `join` to make two strings:

1. "Eats, Shoots, and Leaves" (call this one `with_commas`)
2. "Eats Shoots and Leaves" (call this one `without_commas`)

*Hint:* If you're not sure what `join` does, first try just calling, for example, `"foo".join(book_title_words)`.

BEGIN QUESTION

name: q1\_3

```
[12]: with_commas = ", ".join(book_title_words) #SOLUTION
      without_commas = " ".join(book_title_words) #SOLUTION

      # These lines are provided just to print out your answers.
      print('with_commas:', with_commas)
      print('without_commas:', without_commas)
```

with\_commas: Eats, Shoots, and Leaves

without\_commas: Eats Shoots and Leaves

## 1.2 2. Indexing Arrays

These exercises give you practice accessing individual elements of arrays. In Python (and in many programming languages), elements are accessed by *index*, so the first element is the element at index 0.

**Question 1.** The cell below creates an array of some numbers. Set `third_element` to the third element of `some_numbers`.

BEGIN QUESTION

name: q2\_1

```
[2]: some_numbers = make_array(-1, -3, -6, -10, -15)

      third_element = some_numbers.item(2) #SOLUTION
      third_element
```

[2]: -6

**Question 2.** The next cell creates a table that displays some information about the elements of `some_numbers` and their order. Run the cell to see the partially-completed table, then fill in the missing information (the cells that say "Ellipsis") by assigning `blank_a`, `blank_b`, `blank_c`, and `blank_d` to the correct elements in the table.

BEGIN QUESTION

name: q2\_2

```
[5]: blank_a = "third" #SOLUTION
      blank_b = "fourth" #SOLUTION
      blank_c = 0 #SOLUTION
      blank_d = 3 #SOLUTION
      elements_of_some_numbers = Table().with_columns(
          "English name for position", make_array("first", "second", blank_a,
          ↪blank_b, "fifth"),
          "Index",                      make_array(blank_c, 1, 2, blank_d, 4),
          "Element",                     some_numbers)
```

```
elements_of_some_numbers
```

```
[5]: English name for position | Index | Element
      first                    | 0    | -1
      second                   | 1    | -3
      third                    | 2    | -6
      fourth                   | 3    | -10
      fifth                    | 4    | -15
```

**Question 3.** You'll sometimes want to find the *last* element of an array. Suppose an array has 142 elements. What is the index of its last element?

BEGIN QUESTION

name: q2\_3

```
[10]: index_of_last_element = 142 - 1 # (or just 141) #SOLUTION
```

More often, you don't know the number of elements in an array, its *length*. (For example, it might be a large dataset you found on the Internet.) The function `len` takes a single argument, an array, and returns the *length* of that array (an integer).

**Question 4.** The cell below loads an array called `president_birth_years`. The last element in that array is the most recent birth year of any deceased president. Assign that year to `most_recent_birth_year`.

BEGIN QUESTION

name: q2\_4

```
[12]: president_birth_years = Table.read_table("president_births.csv").column('Birth_
      ↪Year')

      most_recent_birth_year = president_birth_years.
      ↪item(len(president_birth_years)-1) #SOLUTION
      most_recent_birth_year
```

```
[12]: 1917
```

**Question 5.** Finally, assign `sum_of_birth_years` to the sum of the first, tenth, and last birth year in `president_birth_years`.

BEGIN QUESTION

name: q2\_5

```
[14]: sum_of_birth_years = sum(make_array(president_birth_years.item(0),
      ↪president_birth_years.item(9), president_birth_years.
      ↪item(len(president_birth_years) - 1))) #SOLUTION
```

### 1.3 3. Basic Array Arithmetic

**Question 1.** Multiply the numbers 42, 4224, 42422424, and -250 by 157. Assign each variable below such that `first_product` is assigned to the result of `42 * 157`, `second_product` is assigned to the result of `4224 * 157`, and so on.

For this question, **don't** use arrays.

BEGIN QUESTION

name: q3\_1

```
[2]: first_product = 42 * 157 #SOLUTION
     second_product = 4224 * 157 #SOLUTION
     third_product = 42422424 * 157 #SOLUTION
     fourth_product = -250 * 157 #SOLUTION
     print(first_product, second_product, third_product, fourth_product)
```

```
6594 663168 6660320568 -39250
```

**Question 2.** Now, do the same calculation, but using an array called `numbers` and only a single multiplication (`*`) operator. Store the 4 results in an array named `products`.

BEGIN QUESTION

name: q3\_2

```
[4]: numbers = make_array(42, 4224, 42422424, -250) #SOLUTION
     products = numbers * 157 #SOLUTION
     products
```

```
[4]: array([ 6594, 663168, 6660320568, -39250])
```

**Question 3.** Oops, we made a typo! Instead of 157, we wanted to multiply each number by 1577. Compute the correct products in the cell below using array arithmetic. Notice that your job is really easy if you previously defined an array containing the 4 numbers.

BEGIN QUESTION

name: q3\_3

```
[6]: correct_products = numbers * 1577 # SOLUTION
     correct_products
```

```
[6]: array([ 66234, 6661248, 66900162648, -394250])
```

**Question 4.** We've loaded an array of temperatures in the next cell. Each number is the highest temperature observed on a day at a climate observation station, mostly from the US. Since they're from the US government agency [NOAA](#), all the temperatures are in Fahrenheit. Convert them all to Celsius by first subtracting 32 from them, then multiplying the results by  $\frac{5}{9}$ . Make sure to **ROUND** the final result after converting to Celsius to the nearest integer using the `np.round` function.

BEGIN QUESTION

name: q3\_4

```
[8]: max_temperatures = Table.read_table("temperatures.csv").column("Daily Max_
    ↳Temperature")

celsius_max_temperatures = np.round((max_temperatures - 32) * 5/9) #SOLUTION
celsius_max_temperatures
```

```
[8]: array([-4., 31., 32., ..., 17., 23., 16.] )
```

**Question 5.** The cell below loads all the *lowest* temperatures from each day (in Fahrenheit). Compute the size of the daily temperature range for each day. That is, compute the difference between each daily maximum temperature and the corresponding daily minimum temperature. **Give your answer in Celsius!** Make sure **NOT** to round your answer for this question!

BEGIN QUESTION

name: q3\_5

```
[13]: min_temperatures = Table.read_table("temperatures.csv").column("Daily Min_
    ↳Temperature")

celsius_temperature_ranges = (5/9) * (max_temperatures - min_temperatures)_
    ↳#SOLUTION
celsius_temperature_ranges
```

```
[13]: array([ 6.66666667, 10.          , 12.22222222, ..., 17.22222222,
    11.66666667, 11.11111111])
```

## 1.4 4. World Population

Remember that the tests from this point on will **not** necessarily tell you whether or not your answers are correct.

The cell below loads a table of estimates of the world population for different years, starting in 1950. The estimates come from the [US Census Bureau website](#).

```
[2]: world = Table.read_table("world_population.csv").select('Year', 'Population')
world.show(4)
```

<IPython.core.display.HTML object>

The name `population` is assigned to an array of population estimates.

```
[3]: population = world.column(1)
population
```

```
[3]: array([2557628654, 2594939877, 2636772306, 2682053389, 2730228104,
    2782098943, 2835299673, 2891349717, 2948137248, 3000716593,
    3043001508, 3083966929, 3140093217, 3209827882, 3281201306,
```

```

3350425793, 3420677923, 3490333715, 3562313822, 3637159050,
3712697742, 3790326948, 3866568653, 3942096442, 4016608813,
4089083233, 4160185010, 4232084578, 4304105753, 4379013942,
4451362735, 4534410125, 4614566561, 4695736743, 4774569391,
4856462699, 4940571232, 5027200492, 5114557167, 5201440110,
5288955934, 5371585922, 5456136278, 5538268316, 5618682132,
5699202985, 5779440593, 5857972543, 5935213248, 6012074922,
6088571383, 6165219247, 6242016348, 6318590956, 6395699509,
6473044732, 6551263534, 6629913759, 6709049780, 6788214394,
6866332358, 6944055583, 7022349283, 7101027895, 7178722893,
7256490011])

```

In this question, you will apply some built-in Numpy functions to this array. Numpy is a module that is often used in Data Science!

The difference function `np.diff` subtracts each element in an array from the element after it within the array. As a result, the length of the array `np.diff` returns will always be one less than the length of the input array.

The cumulative sum function `np.cumsum` outputs an array of partial sums. For example, the third element in the output array corresponds to the sum of the first, second, and third elements.

**Question 1.** Very often in data science, we are interested understanding how values change with time. Use `np.diff` and `np.max` (or just `max`) to calculate the largest annual change in population between any two consecutive years.

BEGIN QUESTION

name: q4\_1

```

[4]: largest_population_change = max(np.diff(population)) #SOLUTION
      largest_population_change

```

[4]: 87515824

**Question 2.** What do the values in the resulting array represent (choose one)?

```

[6]: np.cumsum(np.diff(population))

```

```

[6]: array([ 37311223,  79143652, 124424735, 172599450, 224470289,
            277671019, 333721063, 390508594, 443087939, 485372854,
            526338275, 582464563, 652199228, 723572652, 792797139,
            863049269, 932705061, 1004685168, 1079530396, 1155069088,
            1232698294, 1308939999, 1384467788, 1458980159, 1531454579,
            1602556356, 1674455924, 1746477099, 1821385288, 1893734081,
            1976781471, 2056937907, 2138108089, 2216940737, 2298834045,
            2382942578, 2469571838, 2556928513, 2643811456, 2731327280,
            2813957268, 2898507624, 2980639662, 3061053478, 3141574331,
            3221811939, 3300343889, 3377584594, 3454446268, 3530942729,
            3607590593, 3684387694, 3760962302, 3838070855, 3915416078,
            3993634880, 4072285105, 4151421126, 4230585740, 4308703704,

```

```
4386426929, 4464720629, 4543399241, 4621094239, 4698861357])
```

- 1) The total population change between consecutive years, starting at 1951.
- 2) The total population change between 1950 and each later year, starting at 1951.
- 3) The total population change between 1950 and each later year, starting inclusively at 1950.

BEGIN QUESTION

name: q4\_2

```
[7]: # Assign cumulative_sum_answer to 1, 2, or 3
      cumulative_sum_answer = 2 # SOLUTION
```

## 1.5 5. Old Faithful

Old Faithful is a geyser in Yellowstone that erupts every 44 to 125 minutes (according to [Wikipedia](#)). People are often told that the geyser erupts every hour, but in fact the waiting time between eruptions is more variable. Let's take a look.

**Question 1.** The first line below assigns `waiting_times` to an array of 272 consecutive waiting times between eruptions, taken from a classic 1938 dataset. Assign the names `shortest`, `longest`, and `average` so that the `print` statement is correct.

BEGIN QUESTION

name: q5\_1

```
[2]: waiting_times = Table.read_table('old_faithful.csv').column('waiting')

      shortest = min(waiting_times) # SOLUTION
      longest = max(waiting_times) # SOLUTION
      average = np.mean(waiting_times) # SOLUTION

      print("Old Faithful erupts every", shortest, "to", longest, "minutes and
      ↳every", average, "minutes on average.")
```

Old Faithful erupts every 43 to 96 minutes and every 70.8970588235294 minutes on average.

**Question 2.** Assign `biggest_decrease` to the biggest decrease in waiting time between two consecutive eruptions. For example, the third eruption occurred after 74 minutes and the fourth after 62 minutes, so the decrease in waiting time was  $74 - 62 = 12$  minutes.

*Hint 1:* You'll need an array arithmetic function mentioned in the textbook. You have also seen this function earlier in the homework!

*Hint 2:* We want to return the absolute value of the biggest decrease.

BEGIN QUESTION

name: q5\_2



```
[9]: biggest_decrease = abs(min((np.diff(waiting_times)))) # SOLUTION
      biggest_decrease
```

```
[9]: 45
```

**Question 3.** If you expected Old Faithful to erupt every hour, you would expect to wait a total of  $60 * k$  minutes to see  $k$  eruptions. Set `difference_from_expected` to an array with 272 elements, where the element at index  $i$  is the absolute difference between the expected and actual total amount of waiting time to see the first  $i+1$  eruptions.

*Hint:* You'll need to compare a cumulative sum to a range. You'll go through `np.arange` more thoroughly in Lab 3, but you can read about it in this [textbook section](#).

For example, since the first three waiting times are 79, 54, and 74, the total waiting time for 3 eruptions is  $79 + 54 + 74 = 207$ . The expected waiting time for 3 eruptions is  $60 * 3 = 180$ . Therefore, `difference_from_expected.item(2)` should be  $|207 - 180| = 27$ .

BEGIN QUESTION

name: q5\_3

```
[13]: difference_from_expected = np.abs(np.cumsum(waiting_times) - np.
      ↪ arange(1,273)*60) # SOLUTION
      difference_from_expected
```

```
[13]: array([ 19,  13,  27,  29,  54,  49,  77, 102,  93, 118, 112,
          136, 154, 141, 164, 156, 158, 182, 174, 193, 184, 171,
          189, 198, 212, 235, 230, 246, 264, 283, 296, 313, 319,
          339, 353, 345, 333, 353, 352, 382, 402, 400, 424, 422,
          435, 458, 462, 455, 477, 476, 491, 521, 515, 535, 529,
          552, 563, 567, 584, 605, 604, 628, 616, 638, 638, 670,
          688, 706, 711, 724, 746, 742, 761, 772, 774, 790, 790,
          808, 824, 847, 862, 884, 894, 899, 912, 940, 956, 976,
          964, 990, 990, 1020, 1010, 1028, 1031, 1043, 1067, 1082, 1073,
          1095, 1097, 1125, 1114, 1137, 1158, 1145, 1169, 1161, 1187, 1208,
          1223, 1222, 1251, 1270, 1269, 1290, 1280, 1305, 1304, 1331, 1324,
          1333, 1350, 1346, 1374, 1395, 1380, 1402, 1397, 1427, 1412, 1435,
          1431, 1460, 1446, 1468, 1459, 1485, 1478, 1497, 1518, 1518, 1540,
          1557, 1573, 1572, 1592, 1581, 1617, 1610, 1627, 1644, 1649, 1670,
          1681, 1691, 1712, 1745, 1738, 1767, 1752, 1778, 1776, 1794, 1800,
          1816, 1819, 1847, 1839, 1872, 1861, 1858, 1875, 1883, 1904, 1925,
          1938, 1928, 1953, 1967, 1962, 1979, 2002, 2025, 2016, 2034, 2058,
          2044, 2067, 2062, 2083, 2080, 2096, 2120, 2137, 2158, 2185, 2202,
          2193, 2211, 2211, 2233, 2264, 2257, 2275, 2261, 2278, 2302, 2291,
          2314, 2325, 2345, 2334, 2349, 2353, 2369, 2362, 2396, 2391, 2407,
          2397, 2419, 2413, 2428, 2446, 2465, 2483, 2501, 2511, 2530, 2540,
          2534, 2560, 2550, 2580, 2574, 2568, 2585, 2604, 2608, 2623, 2610,
          2636, 2639, 2664, 2686, 2683, 2705, 2712, 2726, 2720, 2743, 2756,
          2769, 2797, 2817, 2828, 2851, 2847, 2866, 2884, 2908, 2906, 2929,
          2912, 2912, 2927, 2948, 2934, 2964, 2950, 2964])
```

**Question 4.** Let's imagine your guess for the next wait time was always just the length of the previous waiting time. If you always guessed the previous waiting time, how big would your error in guessing the waiting times be, on average?

For example, since the first three waiting times are 79, 54, and 74, the average difference between your guess and the actual time for just the second and third eruption would be  $\frac{|79-54|+|54-74|}{2} = 22.5$ .

BEGIN QUESTION

name: q5\_4

```
[16]: average_error = np.average(abs(np.diff(waiting_times))) # SOLUTION
      average_error
```

```
[16]: 20.52029520295203
```

## 1.6 6. Tables

**Question 1.** Suppose you have 4 apples, 3 oranges, and 3 pineapples. (Perhaps you're using Python to solve a high school Algebra problem.) Create a table that contains this information. It should have two columns: **fruit name** and **count**. Assign the new table to the variable **fruits**.

**Note:** Use lower-case and singular words for the name of each fruit, like "apple".

BEGIN QUESTION

name: q6\_1

```
[2]: # Our solution uses 1 statement split over 3 lines.
      fruits = Table().with_columns( #SOLUTION
          "fruit name", make_array("apple", "orange", "pineapple"), #SOLUTION
          "count",      make_array(4,      3,      3)) #SOLUTION
      fruits
```

```
[2]: fruit name | count
      apple    | 4
      orange   | 3
      pineapple | 3
```

**Question 2.** The file `inventory.csv` contains information about the inventory at a fruit stand. Each row represents the contents of one box of fruit. Load it as a table named **inventory** using the `Table.read_table()` function. `Table.read_table(...)` takes one argument (a path to a data file in string format) and returns a table.

BEGIN QUESTION

name: q6\_2

```
[4]: inventory = Table.read_table("inventory.csv") #SOLUTION
      inventory
```

```
[4]: box ID | fruit name | count
      53686 | kiwi       | 45
      57181 | strawberry | 123
      25274 | apple     | 20
      48800 | orange    | 35
      26187 | strawberry | 255
      57930 | grape     | 517
      52357 | strawberry | 102
      43566 | peach     | 40
```

**Question 3.** Does each box at the fruit stand contain a different fruit? Set `all_different` to `True` if each box contains a different fruit or to `False` if multiple boxes contain the same fruit

BEGIN QUESTION

name: q6\_3

```
[6]: all_different = False #SOLUTION
      all_different
```

```
[6]: False
```

**Question 4.** The file `sales.csv` contains the number of fruit sold from each box last Saturday. It has an extra column called "price per fruit (\$)" that's the price *per item of fruit* for fruit in that box. The rows are in the same order as the `inventory` table. Load these data into a table called `sales`.

BEGIN QUESTION

name: q6\_4

```
[9]: sales = Table.read_table("sales.csv") #SOLUTION
      sales
```

```
[9]: box ID | fruit name | count sold | price per fruit ($)
      53686 | kiwi       | 3          | 0.5
      57181 | strawberry | 101        | 0.2
      25274 | apple     | 0          | 0.8
      48800 | orange    | 35         | 0.6
      26187 | strawberry | 25         | 0.15
      57930 | grape     | 355        | 0.06
      52357 | strawberry | 102        | 0.25
      43566 | peach     | 17         | 0.8
```

```
[10]: # TEST
      sales.sort(0)
```

```
[10]: box ID | fruit name | count sold | price per fruit ($)
      25274 | apple     | 0          | 0.8
      26187 | strawberry | 25         | 0.15
      43566 | peach     | 17         | 0.8
```

48800	orange	35	0.6
52357	strawberry	102	0.25
53686	kiwi	3	0.5
57181	strawberry	101	0.2
57930	grape	355	0.06

**Question 5.** How many fruits did the store sell in total on that day?

BEGIN QUESTION

name: q6\_5

```
[11]: total_fruits_sold = sum(sales.column("count sold")) #SOLUTION
      total_fruits_sold
```

```
[11]: 638
```

**Question 6.** What was the store's total revenue (the total price of all fruits sold) on that day?

*Hint:* If you're stuck, think first about how you would compute the total revenue from just the grape sales.

BEGIN QUESTION

name: q6\_6

```
[14]: total_revenue = sum(sales.column("count sold") * sales.column("price per fruit_
      ↪($"))) #SOLUTION
      total_revenue
```

```
[14]: 106.85
```

**Question 7.** Make a new table called `remaining_inventory`. It should have the same rows and columns as `inventory`, except that the amount of fruit sold from each box should be subtracted from that box's count, so that the "count" is the amount of fruit remaining after Saturday.

BEGIN QUESTION

name: q6\_7

```
[17]: remaining_inventory = Table().with_columns( #SOLUTION
      "box ID", inventory.column("box ID"), #SOLUTION
      "fruit name", inventory.column("fruit name"), #SOLUTION
      "count", inventory.column("count") - sales.column("count sold")) #SOLUTION
      remaining_inventory
```

```
[17]: box ID | fruit name | count
      53686 | kiwi       | 42
      57181 | strawberry | 22
      25274 | apple     | 20
      48800 | orange    | 0
      26187 | strawberry | 230
      57930 | grape     | 162
```

```
52357 | strawberry | 0
43566 | peach      | 23
```

## 1.7 7. Submission

Once you're finished, select "Save and Checkpoint" in the File menu and then execute the `submit` cell below. The result will contain a link that you can use to check that your assignment has been submitted successfully. If you submit more than once before the deadline, we will only grade your final submission. If you mistakenly submit the wrong one, you can head to [okpy.org](https://okpy.org) and flag the correct version. To do so, go to the website, click on this assignment, and find the version you would like to have graded. There should be an option to flag that submission for grading!

```
[ ]: _ = ok.submit()

[ ]: # For your convenience, you can run this cell to run all the tests at once!
import os
print("Running all tests...")
_ = [ok.grade(q[:-3]) for q in os.listdir("tests") if q.startswith('q') and
    ↪len(q) <= 10]
print("Finished running all tests.")
```